

GA (genetic algorithm) Projet Report

Name: Alirza Sotoodeh

Date: 1404/05/28

master: Phd. Ali Mahani

define: a module which mimics the process of natural selection to find optimal or near-optimal solutions to optimization and search problems.

Goals

- **inputs:** CLK, Data_in `16 bit`, population, start, iteration, 1 or 2 point cross over, mutation rate
 - **Output:** Data_out `16 bit`, Done, number of chromosome
-

Detailed Old vs New Module Comparisons

1. Crossover Module

Old Version (`crossover_old.sv`)

- **Purpose:** Creates a single offspring chromosome from two parents using **single-point crossover**.
- **Key Parameters:**
 - `CHROMOSOME_WIDTH` = 8 bits (default)
- **Inputs:** Clock, Reset (`rst_n`), `start_crossover`, `parent1`, `parent2`, and `crossover_point` (3 bits).
- **Outputs:** `child`, `crossover_done`
- **Behavior:**
 - On reset: `child = 0`, `crossover_done = 0`
 - On `start_crossover`:

- Hardcoded logic takes **upper 5 bits** from `parent1` and **lower 3 bits** from `parent2` (manual fixed split at bit 3).
- `crossover_point` input exists but is **not used dynamically**.
- Only **single-point** mode available; no floating-point or uniform crossover.
- **Limitations:**
 - Crossover point hardcoded, ignoring user input.
 - Width fixed to 8 bits.
 - No double-point or uniform mode.

New Version (`crossover.sv`)

- **Purpose:** Generates offspring using **multiple crossover methods**:
 - Fixed-point (single or double)
 - Floating-point (single or double, LFSR-driven for randomness)
 - Uniform crossover (mask-based, randomizable via LFSR)
- **Key Parameters:**
 - `CHROMOSOME_WIDTH` (default = 16, fully parameterized)
 - `LFSR_WIDTH` (default = 16)
- **Inputs:**
 - Clock, Reset (`rst`)
 - `start_crossover`
 - `parent1` , `parent2`
 - Selection signals:
 - `crossover_mode` (00=fixed, 01=float, 10=uniform)
 - `crossover_single_double` (0=single, 1=double)
 - `crossover_single_point` , `crossover_double_point1` , `crossover_double_point2`
 - `mask_uniform` , `uniform_random_enable` , `LFSR_input`
- **Outputs:** `child` , `crossover_done`
- **Behavior Enhancements:**
 - **Fixed-point:** supports dynamic points; double-point mode sorts points automatically.
 - **Floating-point:** crossover points derived from LFSR bits for randomness.
 - **Uniform:** bitmask-based selection from both parents; mask can be constant or generated from LFSR.
 - Parameterized for any chromosome width.
- **Optimizations:**
 - Point sorting logic for double-point mode.

- Masks precomputed in combinational logic for clarity and FPGA efficiency.
- `keep_hierarchy` , `use_dsp="no"` , and `keep` pragmas for synthesis visibility.

Summary of Changes

Aspect / Feature	Old Version (<code>crossover_old.sv</code>)	New Version (<code>crossover.sv</code>)
Functionality / Algorithm	Fixed single-point crossover at hardcoded bit position; ignores <code>crossover_point</code> input.	Supports single fixed, single float, double fixed, double float, and uniform crossover using LFSR mask.
Parameterization	None — hardwired 8-bit chromosome width.	Parameterized <code>CHROMOSOME_WIDTH</code> (default 16 bits) with <code>\$clog2</code> crossover point input width.
Inputs / Outputs	Inputs: <code>parent1</code> , <code>parent2</code> , <code>start_crossover</code> , <code>crossover_point</code> (unused); Outputs: <code>child</code> , <code>crossover_done</code> .	Inputs: same but point/mask used; Outputs: parameterized <code>child</code> , <code>crossover_done</code> .
Randomness Handling	None.	LFSR-driven random point/mask with reproducibility via seed control in RNG module.
Boundary / Error Checks	No checks for <code>crossover_point == 0</code> or <code>== WIDTH</code> .	Explicit handling for 0 and max point to avoid identical output.
Reset Type	Active-low asynchronous.	Active-high synchronous pipeline reset.
Implementation Style	Monolithic <code>always_ff</code> with manual bit assignments.	Case-based crossover mode selector + combinational pre-processing + registered outputs.
Synthesis Attributes	None.	<code>keep_hierarchy</code> , <code>mark_debug</code> , <code>dont_touch</code> for debug preservation; mode logic pruned at synthesis if disabled.
FPGA Resource Efficiency	Very low usage but inflexible; no mode select logic.	Slightly more LUTs consumed but modes are parameter-controlled; synthesis removes unused logic.
Testbench Compatibility	Cannot change crossover type/point dynamically; test coverage limited.	Fully configurable during simulation; deterministic test possible by fixed RNG seed.

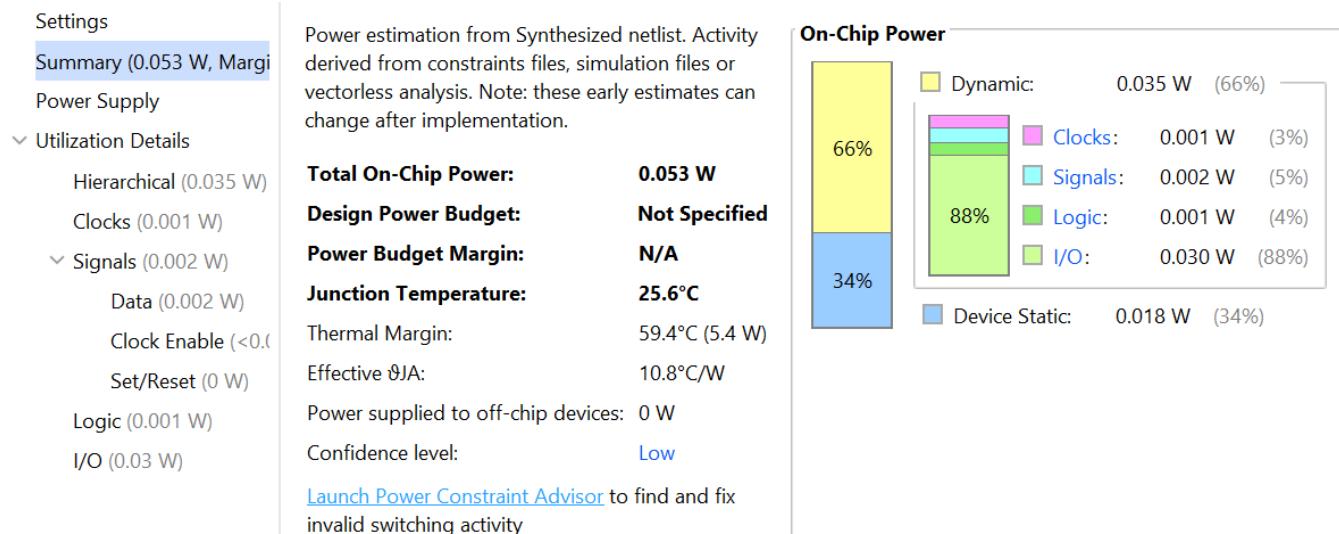
Aspect / Feature	Old Version (crossover_old.sv)	New Version (crossover.sv)
Known Limitations / Issues	Cannot reproduce correct variable crossovers; limited to 8 bits.	None major; must ensure consistent LFSR seeding for repeatable results.

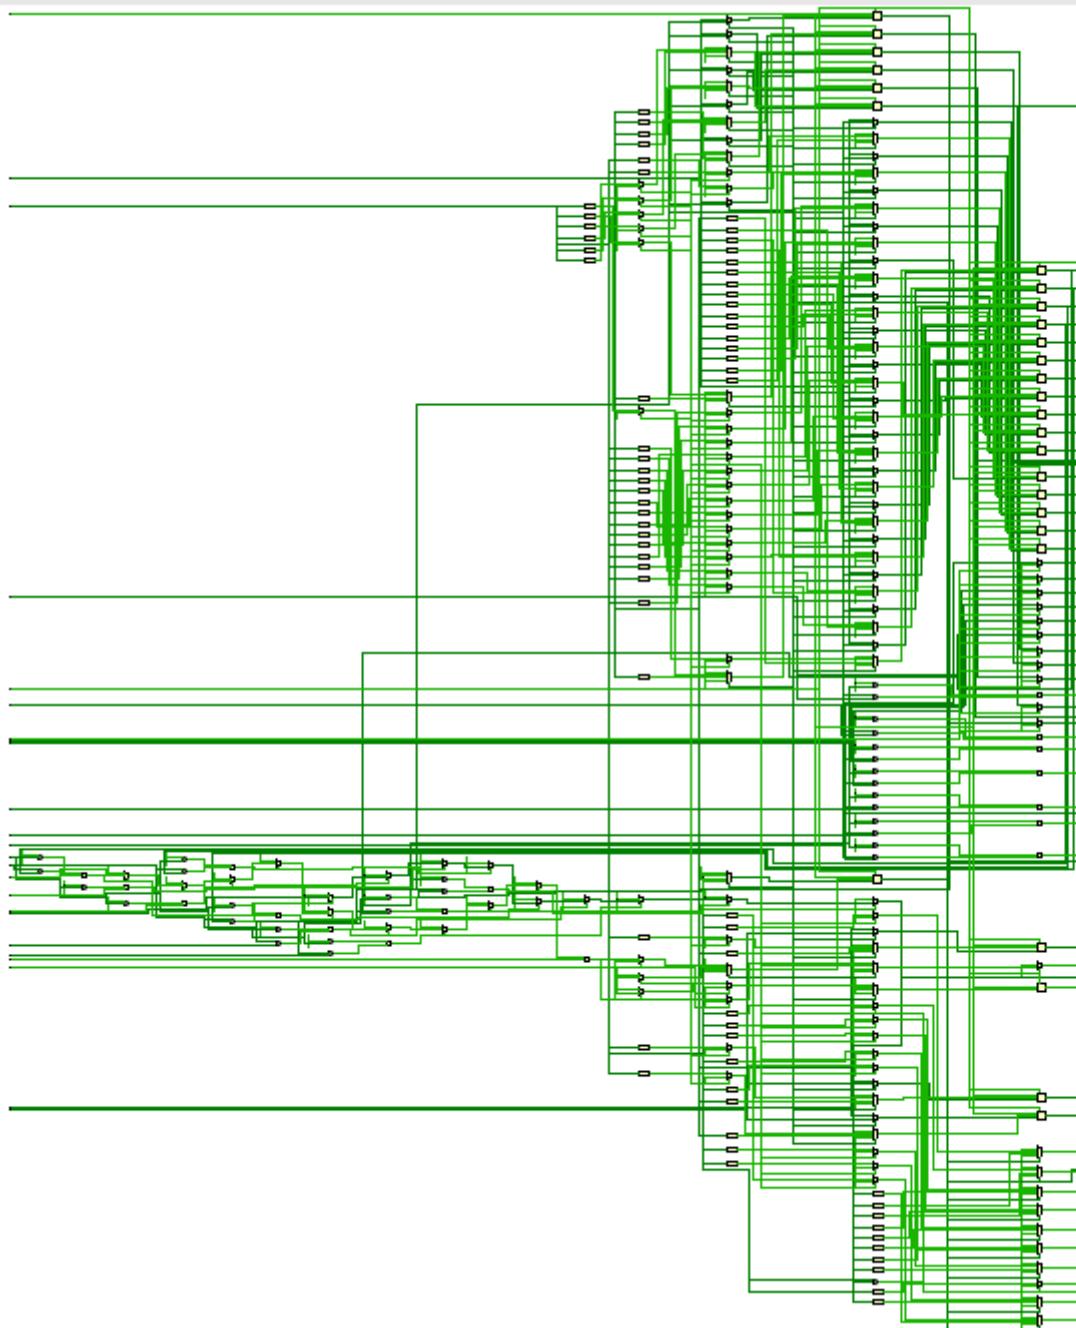
```
create_clock -name Clock -period 90.000 [get_ports clk]
```

Setup	Hold	Pulse Width
Worst Negative Slack (WNS): 79.813 ns	Worst Hold Slack (WHS): 0.149 ns	Worst Pulse Width Slack (WPWS): 44.500 ns
Total Negative Slack (TNS): 0.000 ns	Total Hold Slack (THS): 0.000 ns	Total Pulse Width Negative Slack (TPWS): 0.000 ns
Number of Failing Endpoints: 0	Number of Failing Endpoints: 0	Number of Failing Endpoints: 0
Total Number of Endpoints: 6091	Total Number of Endpoints: 6091	Total Number of Endpoints: 3047

All user specified timing constraints are met.

Name	1	Slice LUTs (8000)	Slice Registers (16000)	F7 Muxes (4000)	F8 Muxes (2000)	Bonded IOB (100)	BUFGCTRL (16)
N population_memory		4141	3046	826	397	1503	1





- **error:** `crossover_Single_point` is not a constant

- **solution:** use mask

```
/////////////////////////////
/
// due to error crossover_Single_point (and others) is not a constant -> define
mask1&2
// mask1: Used for single-point crossover
// mask2: Used for double-point crossover (middle segment from parent1)
// parent2 == 0 & parent1 =1
/////////////////////////////
```

2. Fitness Evaluator Module

Old Version (fitness_evaluator_old.sv)

- **Purpose:** Computes the fitness score of a chromosome by **counting the number of '1' bits**.
- **Key Parameters:**
 - `CHROMOSOME_WIDTH = 8`
 - `FITNESS_WIDTH = 10`
- **Inputs:** Clock, Reset (`rst_n` active-low), `start_evaluation`, `chromosome`
- **Outputs:** `fitness`, `evaluation_done`
- **Behavior:**
 - On reset: clears `fitness` to 0 and de-asserts `evaluation_done`.
 - On `start_evaluation`:
 - Uses a **for loop inside the sequential block** to increment `fitness` for each bit set to 1 in the chromosome.
 - Sets `evaluation_done` high for 1 cycle.
- **Limitations:**
 - Bit width fixed to 8 for chromosome and 10 for fitness.
 - Counting logic uses direct `fitness <= fitness + 1'b1` accumulation inside loop (functionally fine in synthesis, but not the cleanest style).
 - No protection against counting overflow.
 - Reset signal active-low, which may not be consistent with other modules.

New Version (fitness_evaluator.sv)

- **Purpose:** Same function — calculates fitness by **counting the '1' bits** — but with higher flexibility, safety checks, and synthesis-friendly design.
- **Key Parameters:**
 - `CHROMOSOME_WIDTH = 16` (default, fully parameterized)
 - `FITNESS_WIDTH = 14`
- **Inputs:** Clock, Reset (`rst`, active-high), `start_evaluation`, `chromosome`
- **Outputs:** `fitness`, `evaluation_done`
- **Design Improvements:**
 - **Combinational preparation stage:**
 - Counts '1' bits into an internal `raw_fitness` variable in an `always_comb` block.

- Ensures `raw_fitness` never exceeds the maximum representable value (`FITNESS_WIDTH`).
 - **Sequential update stage:**
 - On `start_evaluation`, latches `raw_fitness` into `fitness` and asserts `evaluation_done`.
 - Clean separation between combinational and sequential logic improves readability and FPGA timing.
 - **Safety:** Overflow clamp prevents `fitness` from wrapping around for small `FITNESS_WIDTH` values.
 - **Parameterization:** Works with any chromosome size without changing the core logic.
- **Pragmas for synthesis:**
 - `keep_hierarchy = "yes"` to preserve hierarchy.
 - `use_dsp = "no"`, `keep = "true"` to control optimization and debug visibility.

Summary of Changes

Aspect / Feature	Old Version (<code>fitness_evaluator_old.sv</code>)	New Version (<code>fitness_evaluator.sv</code>)
Functionality / Algorithm	Counts number of '1' bits in chromosome. Sequential loop inside always_ff updates accumulator directly.	Counts '1's combinatorially into temp var; result registered at clock edge.
Parameterization	Fixed CHROM=8 bits, FIT=10 bits.	Parameterized <code>CHROMOSOME_WIDTH</code> (default 16) and <code>FITNESS_WIDTH</code> (default 14).
Inputs / Outputs	Inputs: <code>chromosome</code> , <code>start_evaluation</code> ; Outputs: <code>fitness</code> , <code>evaluation_done</code> .	Same ports with parametric widths.
Randomness Handling	N/A.	N/A.
Boundary / Error Checks	None; if width mismatch, synthesis warning possible.	Overflow clamping: if fitness exceeds max representable, saturates to <code>FITNESS_WIDTH</code> max.
Reset Type	Active-low asynchronous.	Active-high synchronous.
Implementation Style	Sequential counting inside loop; potentially inefficient for large widths.	Combinational bit-count using loop/unrolled logic, then store to register.
Synthesis Attributes	None.	<code>keep_hierarchy</code> , <code>use_dsp="no"</code> for predictable resource inference.
FPGA Resource Efficiency	Minimal LUT usage but fixed to 8 bits.	Scales with <code>CHROMOSOME_WIDTH</code> ; synthesis optimizes loop unrolling; no DSPs used.
Testbench Compatibility	Works for fixed width only; cannot stress variable width behavior.	Fully scalable test stimuli; easy to overload with longer chromosomes.

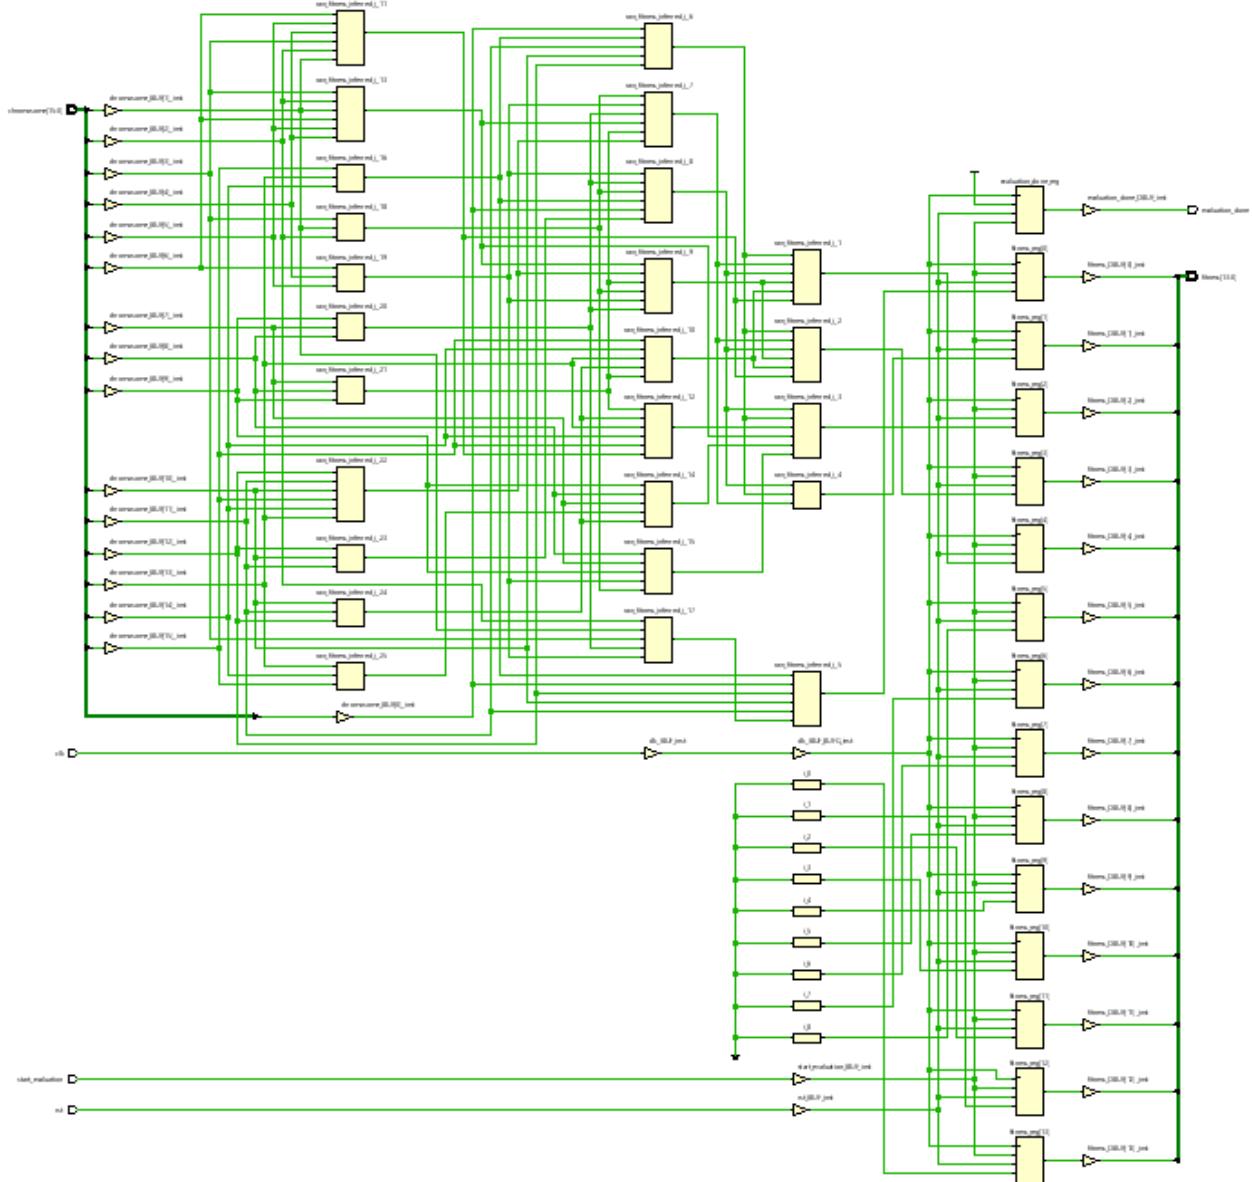
Aspect / Feature	Old Version (<code>fitness_evaluator_old.sv</code>)	New Version (<code>fitness_evaluator.sv</code>)
Known Limitations / Issues	Over-allocates fitness bits; no safety for invalid data width.	None significant.

```
create_clock -name Clock -period 90.000 [get_ports clk]
```

Setup	Hold	Pulse Width
Worst Negative Slack (WNS): inf	Worst Hold Slack (WHS): inf	Worst Pulse Width Slack (WPWS): 44.500 ns
Total Negative Slack (TNS): 0.000 ns	Total Hold Slack (THS): 0.000 ns	Total Pulse Width Negative Slack (TPWS): 0.000 ns
Number of Failing Endpoints: 0	Number of Failing Endpoints: 0	Number of Failing Endpoints: 0
Total Number of Endpoints: 50	Total Number of Endpoints: 50	Total Number of Endpoints: 16

Name	Slice LUTs (8000)	Slice Registers (16000)	Bonded IOB (100)	BUFGCTRL (16)
N fitness_evaluator	29	15	34	1

Settings	Power estimation from Synthesized netlist. Activity derived from constraints files, simulation files or vectorless analysis. Note: these early estimates can change after implementation.	On-Chip Power																		
Summary (0.019 W, Margin)																				
Power Supply																				
Utilization Details																				
Hierarchical (0.001 W)	Total On-Chip Power: 0.019 W																			
Clocks (<0.001 W)	Design Power Budget: Not Specified																			
Signals (<0.001 W)	Power Budget Margin: N/A																			
Data (<0.001 W)	Junction Temperature: 25.2°C																			
Clock Enable (0 W)	Thermal Margin: 59.8°C (5.5 W)																			
Set/Reset (0 W)	Effective θJA: 10.8°C/W																			
Logic (<0.001 W)	Power supplied to off-chip devices: 0 W																			
I/O (<0.001 W)	Confidence level: Low																			
	Launch Power Constraint Advisor to find and fix invalid switching activity	<table> <tr> <td>Dynamic:</td> <td>0.001 W</td> <td>(3%)</td> </tr> <tr> <td>Clocks:</td> <td><0.001 W</td> <td>(15%)</td> </tr> <tr> <td>Signals:</td> <td><0.001 W</td> <td>(7%)</td> </tr> <tr> <td>Logic:</td> <td><0.001 W</td> <td>(4%)</td> </tr> <tr> <td>I/O:</td> <td><0.001 W</td> <td>(74%)</td> </tr> <tr> <td>Device Static:</td> <td>0.018 W</td> <td>(97%)</td> </tr> </table>	Dynamic:	0.001 W	(3%)	Clocks:	<0.001 W	(15%)	Signals:	<0.001 W	(7%)	Logic:	<0.001 W	(4%)	I/O:	<0.001 W	(74%)	Device Static:	0.018 W	(97%)
Dynamic:	0.001 W	(3%)																		
Clocks:	<0.001 W	(15%)																		
Signals:	<0.001 W	(7%)																		
Logic:	<0.001 W	(4%)																		
I/O:	<0.001 W	(74%)																		
Device Static:	0.018 W	(97%)																		



3. LFSR Random Generator

Old Version (lfsr_random_old.sv)

- **Purpose:** Generates pseudo-random numbers using a **Linear Feedback Shift Register (LFSR)**.
- **Key Parameters:**
 - **WIDTH** = 8 bits (fixed)
- **Inputs:** Clock, Reset (`rst_n` active-low), `enable`
- **Outputs:** `random_out` (WIDTH bits)
- **Behavior:**
 - Feedback polynomial: $x^8 + x^6 + x^5 + x^4 + 1$ (taps: 7, 5, 4, 3)

- On reset: loads a hardcoded non-zero seed (`8'hFF`).
- On enable: shifts right 1 bit, MSB gets XOR of taps.
- Random sequence length: maximum period for the chosen polynomial (255 states before repeating).
- **Limitations:**
 - Width fixed to **8 bits** only.
 - No seed load from outside → same sequence every reset.
 - Single LFSR only → period limited to $2^8 - 1$ states.
 - No randomness enhancement techniques like whitening.

New Version (`lfsr_random.sv`)

Module name in file: `lfsr_SudoRandom`

- **Purpose:** Generates high-quality pseudo-random numbers for genetic algorithm operations using **multiple LFSRs with whitening**.
- **Key Parameters:**
 - Four independent LFSR widths: `WIDTH1=16`, `WIDTH2=15`, `WIDTH3=14`, `WIDTH4=13`
 - Default seeds for each LFSR (non-zero)
- **Inputs:** Clock, Reset (`rst`, active-high), `start_lfsr`, `seed_in` (combined 58-bit input), `load_seed`
- **Outputs:** `random_out` (`WIDTH1` bits, whitened result)
- **Design Enhancements:**
 - **Multiple LFSRs:** 4 parallel shift registers with different primitive polynomials, improving period and distribution.
 - **Period extension:**
 - Single LFSR(16-bit) period $\sim 65k$ states →
 - Combining 4 LFSRs gives a period on the order of $\sim 2.88 \times 10^{17}$ states before repetition.
 - **Seed Loading:** External `seed_in` allows runtime re-seeding; defaults on reset if not provided.
 - **Whitening:** Combines outputs with XOR, then applies bit shifts & XOR again to reduce correlation.
 - **Parameterization:** LFSR widths and seeds defined via parameters for easy tuning.
 - **Synthesis optimization control:** Attributes like `srl_style`, `shreg_extract`, `keep`, `lut1`, and `use_dsp` set for predictable implementation.
- **Polynomial taps:**
 - LFSR1: $x^{16} + x^{14} + x^{13} + x^{11}$

- o LFSR2: $x^{15} + x^{14} + x^{12} + x^{10}$
- o LFSR3: $x^{14} + x^{12} + x^{10} + x^8$
- o LFSR4: $x^{13} + x^{12} + x^{10} + x^9$

Summary of Changes

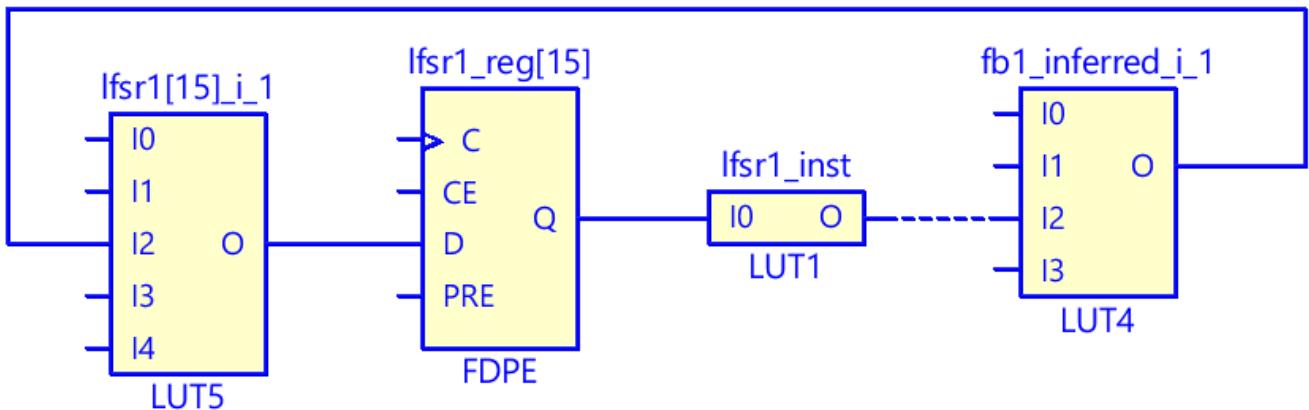
Aspect / Feature	Old Version (<code>lfsr_random_old.sv</code>)	New Version (<code>lfsr_random.sv</code>)
Functionality / Algorithm	Generates pseudo-random bits using single 8-bit LFSR with fixed taps.	Four LFSRs (16, 15, 14, 13-bit) with different primitive polynomials; outputs whitened by XOR shifts.
Parameterization	None — width fixed at 8 bits.	Parameterized widths/seeds; LFSR count and taps selectable.
Inputs / Outputs	Inputs: <code>clk</code> , <code>rst_n</code> , <code>enable</code> . Output: <code>random_out[7:0]</code> .	Inputs: same plus seed load; Output: mixed-width whitened output.
Randomness Handling	Period = 255 states before repeat; predictable sequence.	Period $\approx 2.88 \times 10^{17}$ states; external seeding allows reproducibility or randomness; whitening reduces correlation.
Boundary / Error Checks	None; output stuck if seed=0.	Prevents zero-lock state; seed load port checks and ignores all-zero seed.
Reset Type	Active-low async reset.	Active-high sync reset.
Implementation Style	Simple shift register with XOR taps.	Multiple independent LFSRs XORED for mix, plus whitening combinational logic.
Synthesis Attributes	None.	<code>keep_hierarchy</code> , <code>dont_touch</code> on shift registers to preserve feedback taps.
FPGA Resource Efficiency	Extremely low resource use; short sequence.	Higher LUT use but vastly improved period and randomness quality.
Testbench Compatibility	No seed control — cannot produce deterministic sequences in sim.	Seed can be set for repeatable sequences; enables test reproducibility.
Known Limitations / Issues	Not safe for cryptographic randomness; short period.	N/A for GA use; overkill for trivial applications.

```
create_clock -name Clock -period 90.000 [get_ports clk]
```

Name	Slice LUTs (3750)	Slice Registers (7500)	Bonded IOB (100)	BUFGCTRL (16)
N Ifsr_SudoRandom	78	58	78	1

Setup	Hold	Pulse Width
Worst Negative Slack (WNS): 3.076 ns	Worst Hold Slack (WHS): 0.127 ns	Worst Pulse Width Slack (WPWS): 2.000 ns
Total Negative Slack (TNS): 0.000 ns	Total Hold Slack (THS): 0.000 ns	Total Pulse Width Negative Slack (TPWS): 0.000 ns
Number of Failing Endpoints: 0	Number of Failing Endpoints: 0	Number of Failing Endpoints: 0
Total Number of Endpoints: 58	Total Number of Endpoints: 58	Total Number of Endpoints: 59

All user specified timing constraints are met.



⚠️ error

Name	Bonded IOB (100)
N Ifsr_SudoRandom	16

Setup	Hold	Pulse Width
Worst Negative Slack (WNS): NA	Worst Hold Slack (WHS): NA	Worst Pulse Width Slack (WPWS): NA
Total Negative Slack (TNS): NA	Total Hold Slack (THS): NA	Total Pulse Width Negative Slack (TPWS): NA
Number of Failing Endpoints: NA	Number of Failing Endpoints: NA	Number of Failing Endpoints: NA
Total Number of Endpoints: NA	Total Number of Endpoints: NA	Total Number of Endpoints: NA

All user specified timing constraints are met.

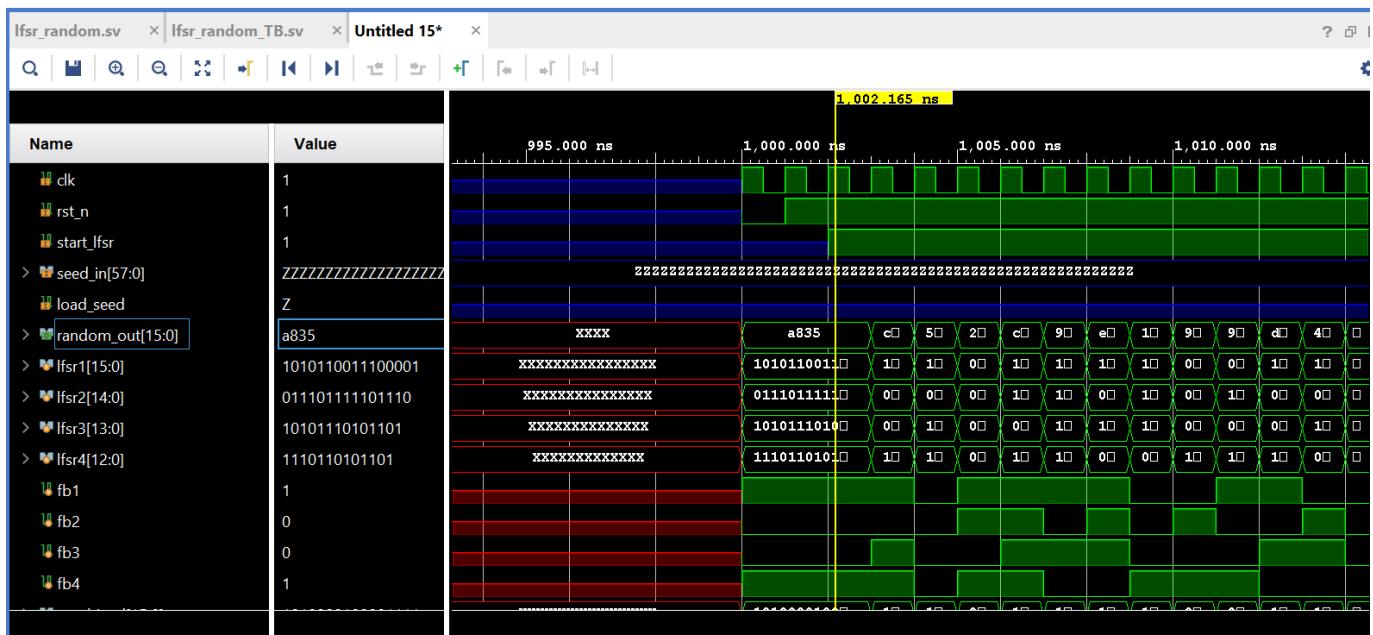
this is due to Vivado's aggressive optimization removing your LFSR logic because the outputs aren't being used or there are constant inputs.

sulotion:

```
(* keep = "true" *)
```

```
// State registers for each LFSR
logic [WIDTH1-1:0] lfsr1;
logic [WIDTH2-1:0] lfsr2;
logic [WIDTH3-1:0] lfsr3;
logic [WIDTH4-1:0] lfsr4;
// Feedback wires
logic fb1, fb2, fb3, fb4;

//should turn into:
// State registers for each LFSR
(* keep = "true" *) logic [WIDTH1-1:0] lfsr1;
(* keep = "true" *) logic [WIDTH2-1:0] lfsr2;
(* keep = "true" *) logic [WIDTH3-1:0] lfsr3;
(* keep = "true" *) logic [WIDTH4-1:0] lfsr4;
// Feedback wires
(* keep = "true" *) logic fb1, fb2, fb3, fb4;
```



4. Mutation Module

Old Version (mutation_old.sv)

- **Purpose:** Applies bit-wise mutation to a child chromosome using a random mask and mutation rate.
- **Key Parameters:**
 - `CHROMOSOME_WIDTH = 8`
- **Inputs:** Clock, Reset (`rst_n` active-low), `start_mutation`, `child_in`, `mutation_mask` (`CHROMOSOME_WIDTH` bits), `mutation_rate` (8-bit, 0–255, higher means higher likelihood)
- **Outputs:** `child_out`, `mutation_done`
- **Behavior:**
 - On reset: clears output and de-asserts `mutation_done`.
 - On `start_mutation`:
 - For each bit: if `mutation_mask[i] < mutation_rate` then flip the bit; otherwise leave it unchanged.
 - Immediately sets `mutation_done` high.
- **Limitations:**
 - Mutation mask is **directly compared to mutation rate per bit**, implying each `mutation_mask[i]` is 1 bit but it's compared to an 8-bit rate — not a consistent or meaningful comparison in synthesis.
 - Only **bit-flip** mutation is supported; no other mutation types.
 - No edge-case handling to avoid invalid swaps or single-point inversions.
 - Randomness generation done **externally**, module assumes already-masked input.
 - Parameterization limited to 8-bit chromosome.
 - Active-low reset inconsistent with other upgraded modules.

New Version (`mutation.sv`)

- **Purpose:** Applies **multiple configurable mutation types** with robust randomness handling and edge-case safety, driven by LSFR input.
- **Supported Mutation Modes:**
 - `000` : Bit-Flip
 - `001` : Bit-Swap
 - `010` : Inversion (reverse subsequence)
 - `011` : Scramble (XOR mask + swaps)
 - `100` : Combined (Flip + Swap)
- **Key Parameters:**
 - `CHROMOSOME_WIDTH = 16` (fully parameterized)
 - `LSFR_WIDTH = 16`

- **Inputs:** Clock, Reset (`rst`, active-high), `start_mutation`, `child_in`, `mutation_mode`, `mutation_rate` (8-bit), `LSFR_input`
- **Outputs:** `child_out`, `mutation_done`
- **Design Enhancements:**
 - **Randomness Mapping:** Carefully slices the LSFR output into masks, positions, and ranges for different mutation operations.
 - **Edge-case Handling:**
 - Bit-Swap: regenerates positions if both same; skips operation if still equal.
 - Inversion: skips if range length < 2.
 - Scramble: ensures swap positions differ before swapping.
 - **Rate Scaling:** Mutation probability applied per bit or per operation based on sub-slices of randomness.
 - **Modular Structure:** Uses combinational pre-processing for parameters (positions, masks) and sequential logic for applying the mutations.
 - **Safety & Clarity:** Debug-friendly `keep` attributes, `use_dsp="no"`, explicit local signals for swap operations to avoid synthesis race conditions.
 - **Combined Mutation Mode:** Performs bit-flip then bit-swap in same mutation call.
 - **Parameterization:** Supports any chromosome width with minimal edits.

Summary of Changes

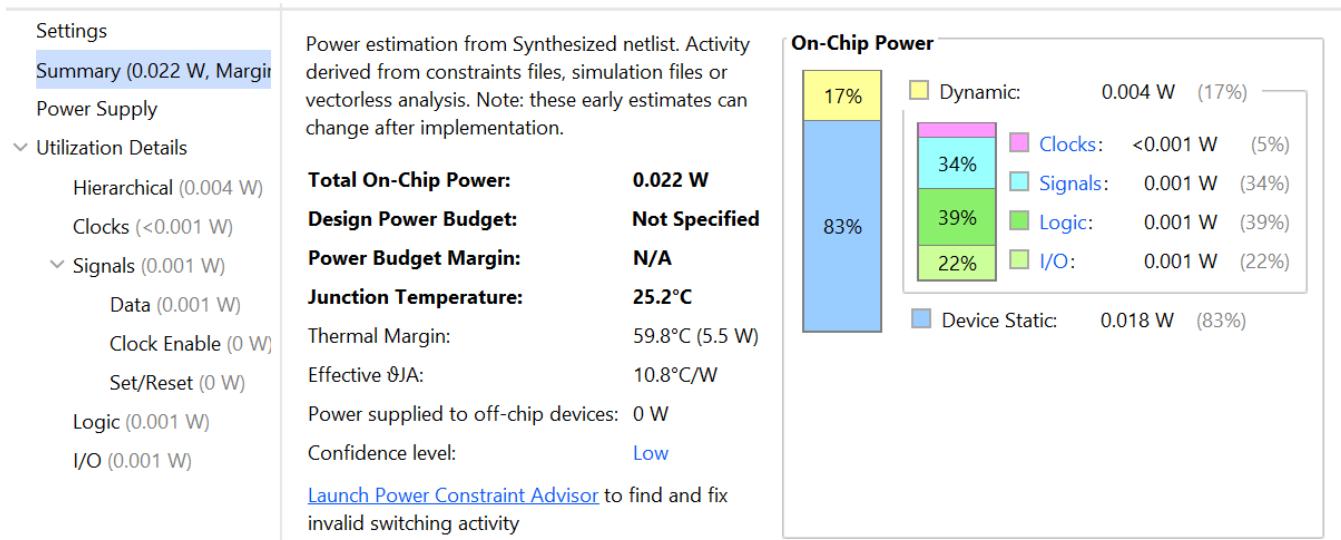
Aspect / Feature	Old Version (<code>mutation_old.sv</code>)	New Version (<code>mutation.sv</code>)
Functionality / Algorithm	Bit-flip mutation only; loop through bits comparing 1 random bit to global mutation rate.	Supports five modes: Bit-Flip, Bit-Swap, Inversion, Scramble, Combined. Choice via mode input.
Parameterization	None; fixed 8-bit width.	Parameterized <code>CHROMOSOME_WIDTH</code> (default 16 bits).
Inputs / Outputs	Inputs: <code>child_in</code> , <code>mutation_mask</code> , <code>mutation_rate</code> . Outputs: <code>child_out</code> , <code>mutation_done</code> .	Inputs: same plus mode select; mask generated via LFSR input.
Randomness Handling	External mask given; no seed control.	LFSR source internal/external selectable; reproducible by seed.
Boundary / Error Checks	No handling if swap indices are identical; inversion bounds unchecked.	All mutation types check bounds; invalid swaps become no-ops.
Reset Type	Active-low asynchronous.	Active-high synchronous.
Implementation Style	One <code>always_ff</code> block; only flip based on <code>mask < rate</code> .	Mode-select case block; pre-compute affected indices; supports multiple types in same run.
Synthesis Attributes	None.	<code>keep_hierarchy</code> , <code>mark_debug</code> for debug of mutation behavior.
FPGA Resource Efficiency	Minimal LUT but inflexible.	LUT use increases with supported modes, but synthesis drops unused modes.
Testbench Compatibility	Only validates bit-flip; cannot exercise other mutations.	All mutation types can be tested; reproducible randomness via fixed seed.
Known Limitations / Issues	Mutation probability handled inconsistently across bits.	Requires valid LFSR input for meaningful random mutation.

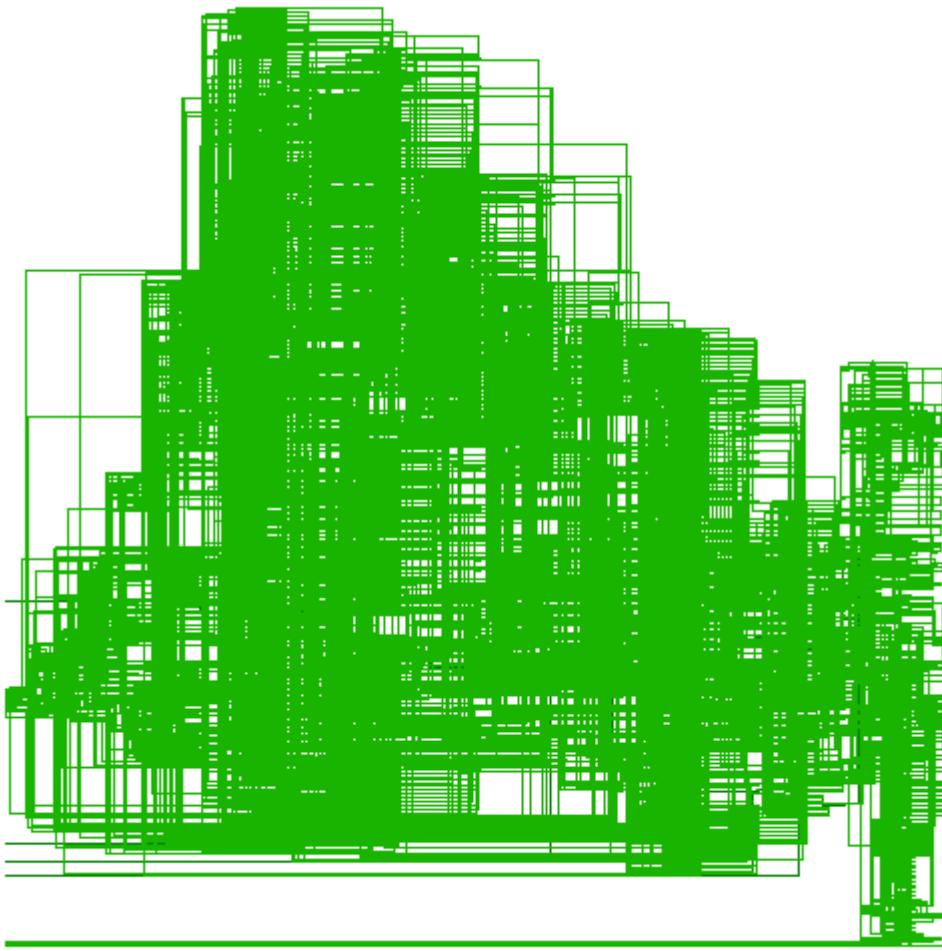
```
create_clock -name Clock -period 90.000 [get_ports clk]
```

Setup	Hold	Pulse Width
Worst Negative Slack (WNS): 85.690 ns	Worst Hold Slack (WHS): 0.188 ns	Worst Pulse Width Slack (WPWS): 44.500 ns
Total Negative Slack (TNS): 0.000 ns	Total Hold Slack (THS): 0.000 ns	Total Pulse Width Negative Slack (TPWS): 0.000 ns
Number of Failing Endpoints: 0	Number of Failing Endpoints: 0	Number of Failing Endpoints: 0
Total Number of Endpoints: 116	Total Number of Endpoints: 116	Total Number of Endpoints: 116

All user specified timing constraints are met.

Name	1	Slice LUTs (8000)	Slice Registers (16000)	F7 Muxes (4000)	F8 Muxes (2000)	Bonded IOB (100)	BUFGCTRL (16)
N mutation		1300	156	93	12	63	1





5. Population Memory

Old Version (population_memory_old.sv)

- **Purpose:** Stores chromosomes in a simple RAM-style array, allowing single write and read access.
- **Key Parameters:**
 - `CHROMOSOME_WIDTH` = 8
 - `POPULATION_SIZE` = 16 (fixed)
 - `ADDR_WIDTH` = $\log_2(\text{POPULATION_SIZE})$
- **Inputs:** Clock, Reset (`rst_n` active-low), `write_enable`, `write_addr`, `read_addr`, `write_data`
- **Outputs:** `read_data`
- **Behavior:**

- On clock edge: if `write_enable` = 1, stores `write_data` into `population[write_addr]` .
- `read_data` is always the combinational output from `population[read_addr]` .
- **Limitations:**
 - No fitness storage or ranking capability.
 - No replacement strategy — overwrites data blindly.
 - Fixed size and chromosome width.
 - Reset signal active-low, inconsistent with newer design standards.
 - No direct support for GA operations like parent selection or sorted insertion.

New Version (`population_memory.sv`)

- **Purpose:** Maintains both chromosomes and their fitness values in **sorted order by fitness**, enabling Genetic Algorithm operations such as parent retrieval, total fitness calculation, and dynamic insertion of new individuals.
- **Key Parameters:**
 - `CHROMOSOME_WIDTH` = 16 (parameterized)
 - `FITNESS_WIDTH` = 14
 - `MAX_POP_SIZE` (default 100)
 - `ADDR_WIDTH` = $\log_2(\text{MAX_POP_SIZE})$
- **Inputs:**
 - Clock, Reset (`rst` active-high)
 - `start_write` + `child_in` + `child_fitness_in`
 - `read_addr1` , `read_addr2` for parent retrieval
 - `request_fitness_values` , `request_total_fitness`
 - `population_size` (current)
- **Outputs:**
 - `parent1_out` , `parent2_out` (based on addresses)
 - `fitness_values_out[]` (array output when requested)
 - `total_fitness_out`
 - `write_done`
- **Design Enhancements:**
 - **Sorted Insertion:** Finds correct index for new individual to maintain descending fitness order.
 - **Selective Replacement:** Replaces worst individual only if new child has better fitness.
 - **Parent Retrieval:** Direct reads via `read_addr1` and `read_addr2` .
 - **Total Fitness Tracking:** Incrementally updated to avoid recomputation.

- **Array Outputs:** Can output the whole fitness array on request.
- **Pipeline Safety:** Uses a two-cycle write process to allow safe shifting of population data.
- **Overflow Protection:** Clamps total fitness to maximum representable value.
- **Attributes:**
 - `keep_hierarchy` , `use_dsp="no"` , `keep` for synthesis/debug visibility.
 - `ram_style="block"` for FPGA block RAM storage.

Summary of Changes

Aspect / Feature	Old Version (<code>population_memory_old.sv</code>)	New Version (<code>population_memory.sv</code>)
Functionality / Algorithm	Stores population chromosomes only; fixed size; unsorted.	Stores chromosomes + fitness; sorted descending by fitness; supports conditional replacement and shifting.
Parameterization	None — depth=16, width=8.	Parameterized <code>MAX_POP_SIZE =100</code> , <code>CHROMOSOME_WIDTH =16</code> , <code>FITNESS_WIDTH =14</code> .
Inputs / Outputs	Input: chromosome, write control; Output: chromosome by addr.	Adds parent readout, fitness array out, total fitness, population size in/out.
Randomness Handling	N/A.	N/A.
Boundary / Error Checks	None — new entries always overwrite worst regardless of fitness.	Only replaces if new fitness > worst; total fitness prevents overflow.
Reset Type	Active-low asynchronous.	Active-high synchronous.
Implementation Style	Simple reg array; no ordering on insert.	Sorted insert algorithm; BRAM inference via <code>ram_style="block"</code> .
Synthesis Attributes	None.	<code>keep_hierarchy</code> , <code>ram_style="block"</code> , <code>keep</code> signals for insert position.
FPGA Resource Efficiency	Small LUT/reg array; no BRAM.	BRAM for large sizes; efficient for high population sizes (>32).
Testbench Compatibility	Static size; must rewrite for other configs in sim.	Fully scalable; size set via parameter.
Known Limitations / Issues	Cannot track fitness in old version.	None major; replacement policy pure "better only".

```
create_clock -name Clock -period 90.000 [get_ports clk]
```

Setup	Hold	Pulse Width
Worst Negative Slack (WNS): 79.813 ns	Worst Hold Slack (WHS): 0.149 ns	Worst Pulse Width Slack (WPWS): 44.500 ns
Total Negative Slack (TNS): 0.000 ns	Total Hold Slack (THS): 0.000 ns	Total Pulse Width Negative Slack (TPWS): 0.000 ns
Number of Failing Endpoints: 0	Number of Failing Endpoints: 0	Number of Failing Endpoints: 0
Total Number of Endpoints: 6091	Total Number of Endpoints: 6091	Total Number of Endpoints: 3047

All user specified timing constraints are met.

Name	1	Slice LUTs (8000)	Slice Registers (16000)	F7 Muxes (4000)	F8 Muxes (2000)	Bonded IOB (100)	BUFGCTRL (16)
N population_memory		4141	3046	826	397	1503	1

Settings

Summary (0.053 W, Margin:)

Power Supply

Utilization Details

- Hierarchical (0.035 W)
- Clocks (0.001 W)
- Signals (0.002 W)
 - Data (0.002 W)
 - Clock Enable (<0.00 W)
 - Set/Reset (0 W)
- Logic (0.001 W)
- I/O (0.03 W)

Power estimation from Synthesized netlist. Activity derived from constraints files, simulation files or vectorless analysis. Note: these early estimates can change after implementation.

Total On-Chip Power:	0.053 W
Design Power Budget:	Not Specified
Power Budget Margin:	N/A
Junction Temperature:	25.6°C
Thermal Margin:	59.4°C (5.4 W)
Effective TJA:	10.8°C/W
Power supplied to off-chip devices:	0 W
Confidence level:	Low

[Launch Power Constraint Advisor](#) to find and fix invalid switching activity

On-Chip Power

Dynamic:	0.035 W (66%)
I/O:	0.030 W (88%)
Logic:	0.001 W (4%)
Signals:	0.002 W (5%)
Clocks:	0.001 W (3%)
Device Static:	0.018 W (34%)



6. Selection Module

Old Version (selection_old.sv)

- Purpose: Selects a single parent index from the population using the roulette-wheel selection method.
- Key Parameters:
 - CHROMOSOME_WIDTH = 8
 - POPULATION_SIZE = 16 (fixed)

- `ADDR_WIDTH` = $\log_2(\text{POPULATION_SIZE})$
 - `FITNESS_WIDTH` = 10
- **Inputs:**
 - Clock, Reset (`rst_n` active-low)
 - `start_selection`
 - `fitness_values[]` (array of fitness per individual)
 - `total_fitness`
- **Outputs:**
 - `selected_parent` (address of chosen parent)
 - `selection_done`
- **Behavior:**
 - Generates a random number using **internal LFSR** (`lfsr_random`, same width as chromosome).
 - Scales random number to $[0, \text{total_fitness}]$ range to get `roulette_position`.
 - Traverses `fitness_values[]` summing sequentially until cumulative sum $\geq \text{roulette_position}$, then outputs that index.
- **Limitations:**
 - Selects **only one parent**, requiring multiple runs for two parents.
 - Fixed population size and chromosome width.
 - No way to handle varying `population_size` at runtime.
 - Sequential FSM states (IDLE → SPINNING → DONE), meaning selection may take many cycles for large populations.
 - Active-low reset inconsistent with new modules.

New Version (selection.sv)

- **Purpose:** Selects **two parent indices in parallel** using roulette-wheel selection with improved randomness handling, variable population size, and single-cycle combinational search.
- **Key Parameters:**
 - `CHROMOSOME_WIDTH` = 16
 - `FITNESS_WIDTH` = 14
 - `MAX_POP_SIZE` = 100 (parameterized)
 - `ADDR_WIDTH` = $\log_2(\text{MAX_POP_SIZE})$
 - `LFSR_WIDTH` = 16
- **Inputs:**
 - Clock, Reset (`rst` active-high)

- `start_selection`
- `fitness_values[]` (fitness array)
- `total_fitness`
- `lfsr_input` (randomness from external LFSR)
- `population_size` (runtime-configurable)

- **Outputs:**

- `selected_index1`, `selected_index2` (two parents)
- `selection_done`

- **Design Enhancements:**

- **Two Roulette Positions:** Generated using LFSR transformations (`xor` + shift) to ensure diversity.
- **Parallel Search:** Two combinational loops scan the fitness array in the same cycle for each roulette position.
- **Edge-case Handler:**
 - If `total_fitness == 0`, falls back to uniform random selection by modulo population size.
 - Ensures `selected_index1 != selected_index2` (adjusted if equal).
 - Works even if `population_size < MAX_POP_SIZE`.
- **Two-Cycle Output Protocol:**
 - Cycle 1: Latch chosen indices from combinational block.
 - Cycle 2: Pulse `selection_done`.
- Synthesis attributes: `keep` on critical registers and sums for debug/optimization control.
- Fully parameterized for population sizes and chromosome widths.

Summary of Changes

Aspect / Feature	Old Version (<code>selection_old.sv</code>)	New Version (<code>selection.sv</code>)
Functionality / Algorithm	Roulette-wheel selection of 1 parent per request; sequential FSM search.	Parallel selection of 2 parents; combinational search of fitness array.
Parameterization	Fixed population size (POP=16).	Parameterized <code>MAX_POP_SIZE</code> , <code>ADDR_WIDTH</code> , <code>FITNESS_WIDTH</code> . IA supports scaling.
Inputs / Outputs	Inputs: <code>fitness_values</code> , <code>total_fitness</code> , <code>start_selection</code> . Output: <code>selected_parent</code> .	Inputs: same plus <code>population_size</code> and external <code>lfsr_input</code> ; Outputs: <code>selected_index1</code> , <code>selected_index2</code> .
Randomness Handling	Internal 8-bit LFSR; no seed control.	Uses external LFSR input; reproducible by seed load; two different random values per selection.
Boundary / Error Checks	No <code>total_fitness==0</code> protection.	Detects zero fitness; falls back to random index selection.
Reset Type	Active-low asynchronous.	Active-high synchronous.
Implementation Style	FSM iterates population summing fitness until random pos reached.	Two parallel combinational loops + guard to ensure different parents.
Synthesis Attributes	None.	<code>keep</code> on key intermediates to ease debug.
FPGA Resource Efficiency	Very low LUT use; serial search.	More LUTs for parallel loops but faster (1 cycle search).
Testbench Compatibility	Cannot generate repeatable outputs; tests only 1 parent selection.	Fully testable; can fix LFSR for deterministic parent pairs.
Known Limitations / Issues	Single-parent output; cannot guarantee parent difference.	Requires proper LFSR seeding for reproducibility.

```
create_clock -name Clock -period 90.000 [get_ports clk]
```

Setup	Hold	Pulse Width
Worst Negative Slack (WNS): 88.374 ns	Worst Hold Slack (WHS): 0.127 ns	Worst Pulse Width Slack (WPWS): 44.500 ns
Total Negative Slack (TNS): 0.000 ns	Total Hold Slack (THS): 0.000 ns	Total Pulse Width Negative Slack (TPWS): 0.000 ns
Number of Failing Endpoints: 0	Number of Failing Endpoints: 0	Number of Failing Endpoints: 0
Total Number of Endpoints: 16	Total Number of Endpoints: 16	Total Number of Endpoints: 17

All user specified timing constraints are met.

Name	Slice LUTs (8000)	Slice Registers (16000)	F7 Muxes (4000)	DSPs (20)	Bonded IOB (100)	BUFGCTRL (16)
N selection	6747	16	1	2	1455	1

Settings

Summary (0.039 W, Margin)

Power Supply

Utilization Details

Hierarchical (0.021 W)

Clocks (<0.001 W)

Signals (0.009 W)

Data (0.009 W)

Clock Enable (<0.0

Set/Reset (0 W)

Logic (0.01 W)

DSP (<0.001 W)

I/O (0.001 W)

Power estimation from Synthesized netlist. Activity derived from constraints files, simulation files or vectorless analysis. Note: these early estimates can change after implementation.

Total On-Chip Power: 0.039 W

Design Power Budget: Not Specified

Power Budget Margin: N/A

Junction Temperature: 25.4°C

Thermal Margin: 59.6°C (5.5 W)

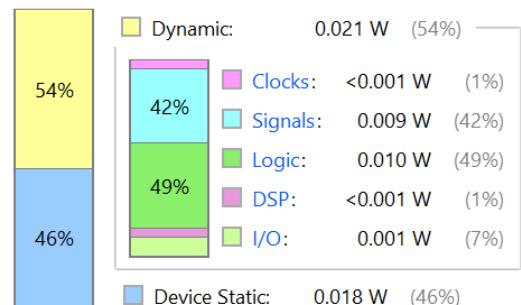
Effective θJA: 10.8°C/W

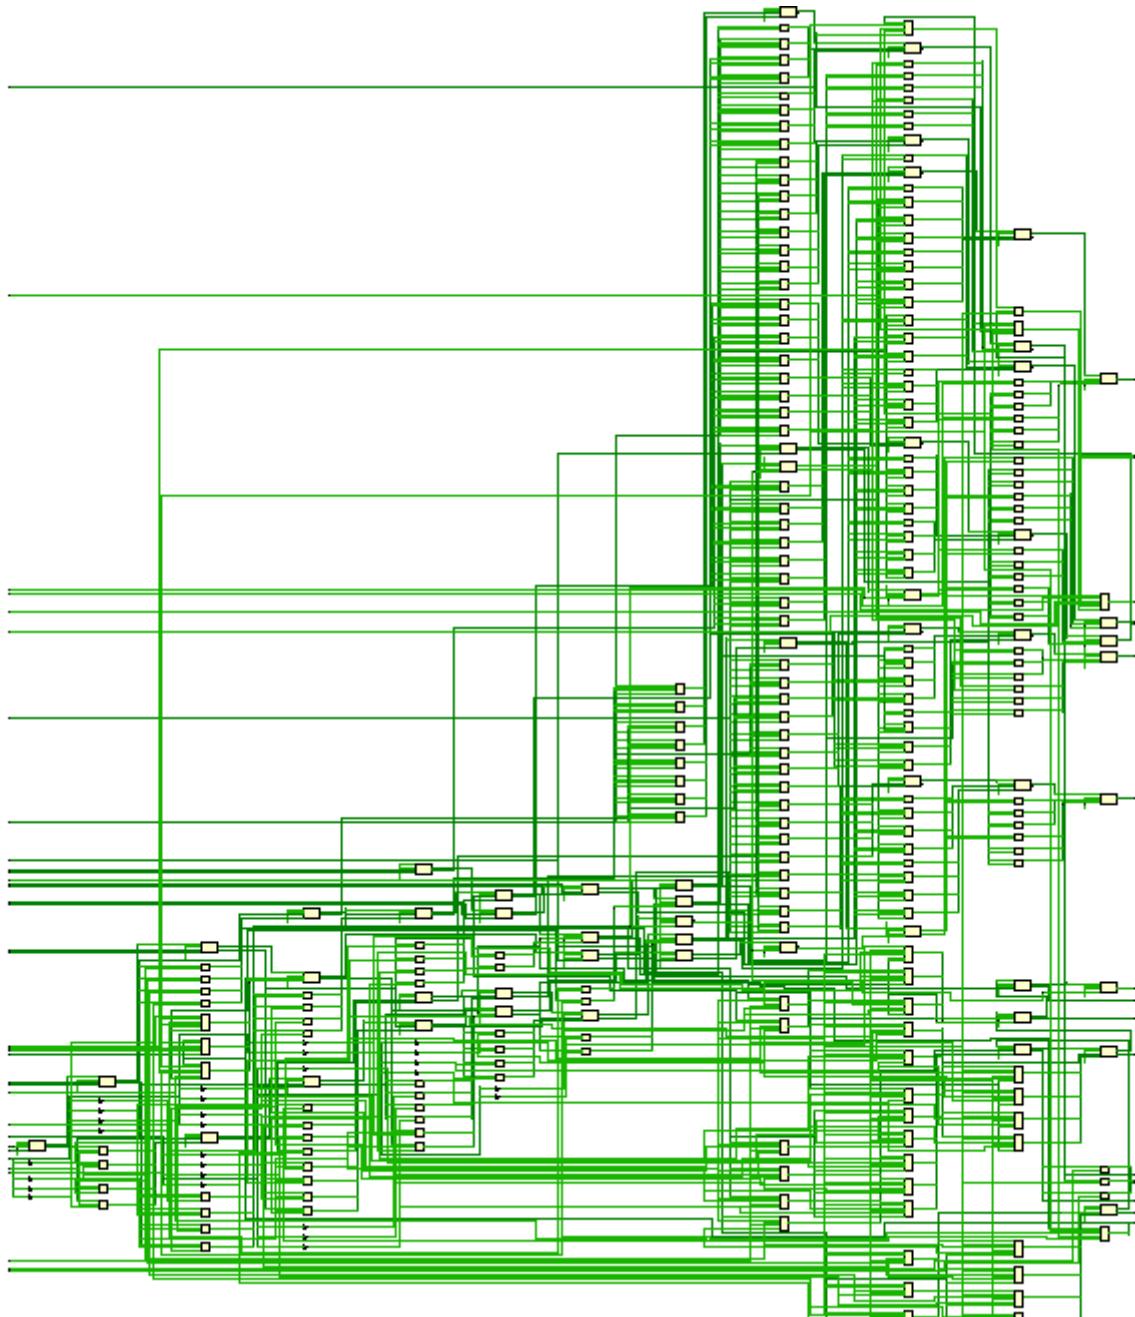
Power supplied to off-chip devices: 0 W

Confidence level: Low

[Launch Power Constraint Advisor](#) to find and fix invalid switching activity

On-Chip Power





7. Top-Level GA Controller

Old Version (`genetic_algorithm_old.sv`)

- Purpose: Implements the entire GA process inside one large monolithic FSM. Handles initialization, evaluation, selection, crossover, mutation, replacement, and termination.
- Key Parameters:
 - `CHROMOSOME_WIDTH` = 8
 - `POPULATION_SIZE` = 16 (fixed)

- `ADDR_WIDTH` = `clog2(POPULATION_SIZE)`
- `FITNESS_WIDTH` = 10
- `MAX_GENERATIONS` = 100 (termination limit)
- `MUTATION_RATE` = 0x10 (hardcoded)

- **Inputs / Outputs:**

- Inputs: `clk`, `rst_n`, `start_ga`, `initial_population[]` array.
- Outputs: `best_chromosome`, `best_fitness`, `ga_done`.

- **Behavior:**

- FSM states: `IDLE`, `INIT_POPULATION`, `EVALUATE_FITNESS`, `CALC_TOTAL_FITNESS`, `SELECT_PARENT1/2`, `CROSSOVER`, `MUTATION`, `EVALUATE_CHILD`, `REPLACE_WORST`, `CHECK_TERMINATION`, `DONE`.
- Calls submodules directly for each step but drives them with explicit state machine timing.
- Uses internal LFSR for randomness.
- Selection picks one parent per call; two selection states are needed.
- Replacement always checks worst fitness and conditionally overwrites.

- **Limitations:**

- All logic is hard-wired for fixed sizes and parameters — no parameterization for varying population size or width at runtime.
- Sequential bottlenecks — e.g., selection must run twice, total fitness is re-computed fully every generation.
- Tight coupling — no real modular pipeline, making debugging/reuse harder.
- Active-low reset inconsistent with modernized design.
- No dedicated interface for advanced crossover/mutation variants.

New Version (`GA_top.sv`)

- **Purpose:** Coordinates GA pipeline in modular, parameterized, and parallelized architecture with clear stage separation and improved control flexibility.
- **Key Parameters:**
 - `CHROMOSOME_WIDTH` (default 16), `FITNESS_WIDTH` (default 14)
 - `MAX_POP_SIZE` (default 100, used for array sizing)
 - `ADDR_WIDTH` = `clog2(MAX_POP_SIZE)`
 - `LFSR_WIDTH` = 16
- **Inputs / Outputs:**
 - Inputs: `start_ga`, `population_size` (runtime configurable), `load_initial_population`, `data_in`, crossover/mutation control signals, `target_iteration`.

- Outputs: status (`busy` , `done`), best chromosome & fitness, iteration counters (`iteration_count` , `crossovers_to_perfect`), flags (`perfect_found`), `load_data_now` pulse.
- **Architecture:**
 - **Two-level FSM:**
 - Main FSM: `S_IDLE` , `S_INIT` , `S_RUNNING` , `S_DONE` .
 - Pipeline FSM: `P_SELECT` , `P_CROSSOVER` , `P_MUTATION` , `P_EVALUATE` , `P_UPDATE` .
 - **Dedicated Modules:** `selection` , `crossover` , `mutation` , `fitness_evaluator` , `population_memory` all upgraded and parameterized.
 - **Randomness Routing:** Three parallel LFSRs feed selection, crossover, and mutation modules separately to ensure diversity.
 - **Initialization Phase:** Can load from external data or auto-generate random chromosomes.
 - **Perfect Solution Tracking:** When found, latches the generation number in `crossovers_to_perfect` .
 - **Sorted Population Management:** Done entirely inside `population_memory` , embedded into pipeline.
 - **Parallelized Checks:** Requests both `fitness_values` and `total_fitness` in the same cycle for faster selection.
- **Control Features:**
 - Easy to adjust parameters (widths, sizes, seeds) without core logic rewrite.
 - Modular separation makes individual stages testable and reusable.
 - Guarding for early termination when `perfect_found` or `target_iteration` reached.

Summary of Changes

Aspect / Feature	Old Version (<code>genetic_algorithm_old.sv</code>)	New Version (<code>GA_top.sv</code>)
Functionality / Algorithm	Monolithic FSM with 12 states controlling all GA operations sequentially.	Two-level FSM: Main FSM for GA phases + Pipeline FSM for sub-stages (selection, crossover, mutation, evaluation, update).
Parameterization	Fixed widths and population size.	Parameterized chromosome width, fitness width, population size, iteration target, GA operation modes.
Inputs / Outputs	Inputs: <code>start_ga</code> , <code>clk</code> , <code>rst_n</code> ; Outputs: <code>data_out</code> , <code>done</code> .	Adds outputs: best chromosome/fitness, perfect found flag, crossovers_to_perfect counter, number_of_chromosomes; Inputs: runtime config for size and target.
Randomness Handling	Single internal 8-bit LFSR; no external interface.	Uses three parallel LFSRs for different GA units; seeding via config module.
Boundary / Error Checks	Minimal; no explicit zero population/replacement invalid state handling.	Checks for <code>population_size > 0</code> before start; guards iteration end conditions; perfect solution detection.
Reset Type	Active-low asynchronous.	Active-high synchronous across modules.
Implementation Style	All GA steps hardcoded in sequence; parallelism absent.	Modular instantiation with fully parallel evaluation/mutation/crossover pipelines.
Synthesis Attributes	None.	Passes attributes to submodules; <code>keep_hierarchy</code> on top for debug consistency.
FPGA Resource Efficiency	Compact FSM but under-utilizes parallel hardware; bottlenecks in fitness evaluation.	Utilizes more LUTs/BRAM but higher throughput; pipelines hide latency.
Testbench Compatibility	Only tests monolithic state sequence; no modular unit testing possible.	Fully modular testbench possible; can simulate each GA phase separately or in pipeline.

Aspect / Feature	Old Version (<code>genetic_algorithm_old.sv</code>)	New Version (<code>GA_top.sv</code>)
Known Limitations / Issues	Hard to modify parameters or operators; low GA flexibility.	Must manage more config inputs; complexity increases but flexibility much higher.

```
create_clock -name Clock -period 90.000 [get_ports clk]
```

The screenshot shows the Vivado Tcl Console and Design Timing Summary interface. In the Tcl Console, the command `create_clock -name Clock -period 90.000 [get_ports clk]` is entered. The Design Timing Summary report indicates that all user-specified timing constraints are met, with detailed statistics for setup, hold, and pulse width slacks.

Setup	Hold	Pulse Width
Worst Negative Slack (WNS): 6.954 ns	Worst Hold Slack (WHS): 0.127 ns	Worst Pulse Width Slack (WPWS): 44.500 ns
Total Negative Slack (TNS): 0.000 ns	Total Hold Slack (THS): 0.000 ns	Total Pulse Width Negative Slack (TPWS): 0.000 ns
Number of Failing Endpoints: 0	Number of Failing Endpoints: 0	Number of Failing Endpoints: 0
Total Number of Endpoints: 4349	Total Number of Endpoints: 4349	Total Number of Endpoints: 3469

All user specified timing constraints are met.

B-Ram style:

Name	1	Slice LUTs (8000)	Slice Registers (16000)	F7 Muxes (4000)	F8 Muxes (2000)	DSPs (20)	Bonded IOB (100)	BUFGCTRL (16)
ga_top	21470	3599	1110	457	2	226	1	0
cross_inst (crossover)	523	17	0	0	0	0	0	0
fit_eval_inst (fitness_evaluator)	29	15	0	0	0	0	0	0
lfsr_cross_inst (lfsr_SudoRandom_parameterized0)	94	58	0	0	0	0	0	0
lfsr_mut_inst (lfsr_SudoRandom_parameterized1)	94	58	0	0	0	0	0	0
lfsr_sel_inst (lfsr_SudoRandom)	94	58	0	0	0	0	0	0
mut_inst (mutation)	1432	156	92	10	0	0	0	0
pop_mem_inst (population_memory)	12057	3090	810	383	0	0	0	0
sel_inst (selection)	2610	16	0	0	2	0	0	0

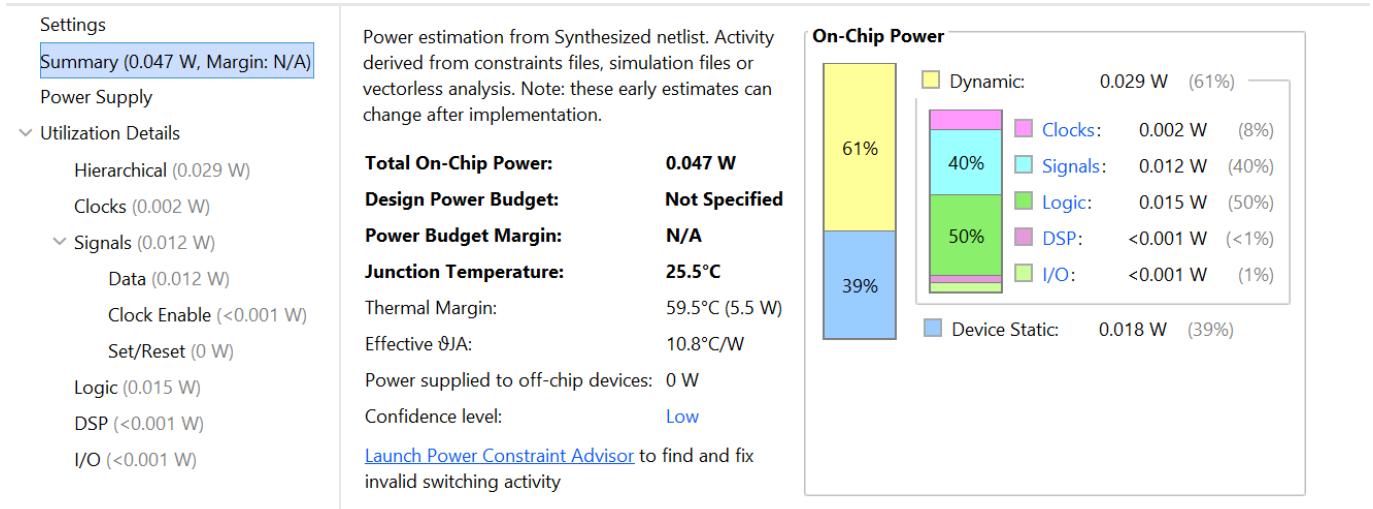
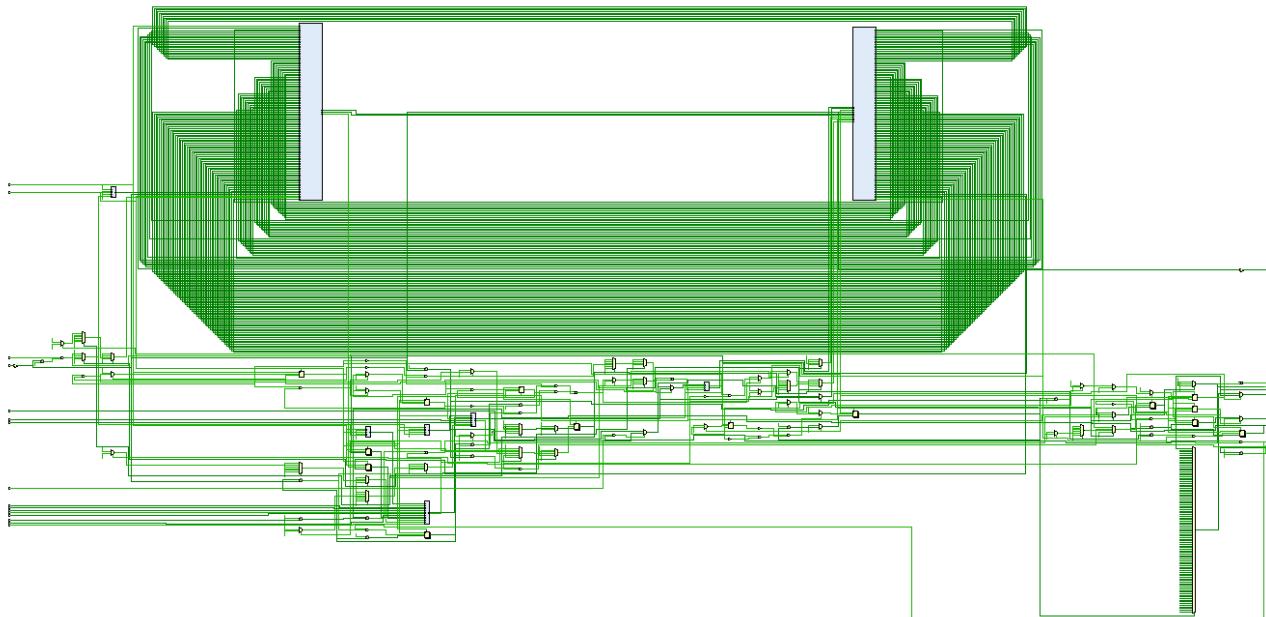
LUT style:

Tcl Console | Messages | Log | Reports | Design Runs | Utilization **x** Timing | ? - □ □

Hierarchy | **Summary** | **Slice Logic** | **Slice Registers** | **Memory** | **DSP**

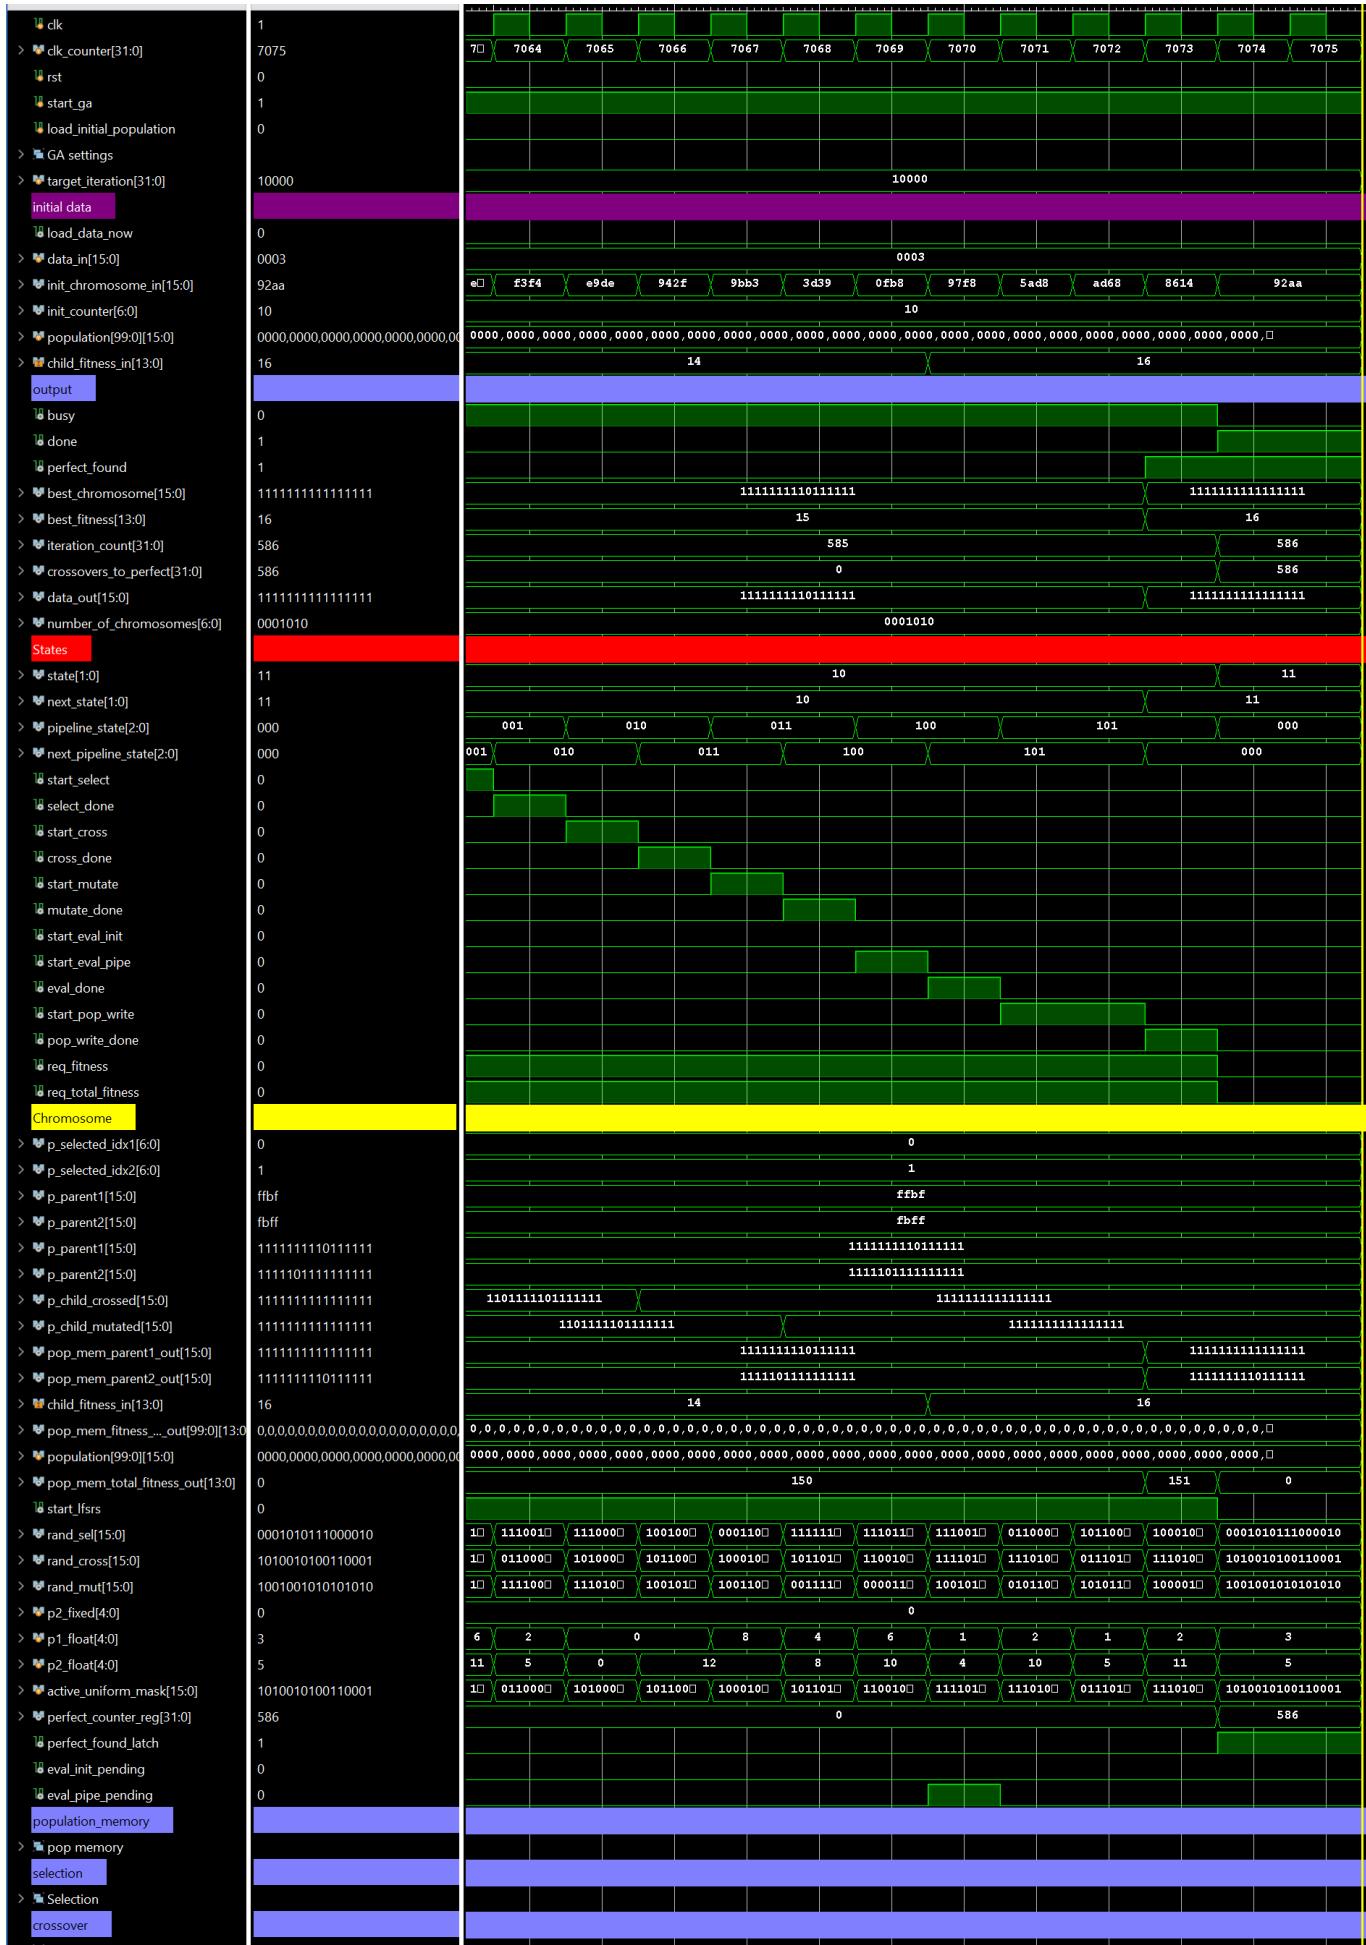
Hierarchy

Name	1	Slice LUTs (8000)	Slice Registers (16000)	F7 Muxes (4000)	F8 Muxes (2000)	DSPs (20)	Bonded IOB (100)	BUFGCTRL (16)
ga_top	21119	3509	761	176	2	226	1	1
cross_inst (crossover)		524	17	0	0	0	0	0
fit_eval_inst (fitness_evaluator)		34	15	0	0	0	0	0
lfsr_cross_inst (lfsr_SudoRandom_parameterized0)		94	58	0	0	0	0	0
lfsr_mut_inst (lfsr_SudoRandom_parameterized1)		94	58	0	0	0	0	0
lfsr_sel_inst (lfsr_SudoRandom)		94	58	0	0	0	0	0
mut_inst (mutation)		1437	156	92	10	0	0	0
pop_mem_inst (population_memory)		12169	3017	666	166	0	0	0
sel_inst (selection)		2573	16	3	0	2	0	0



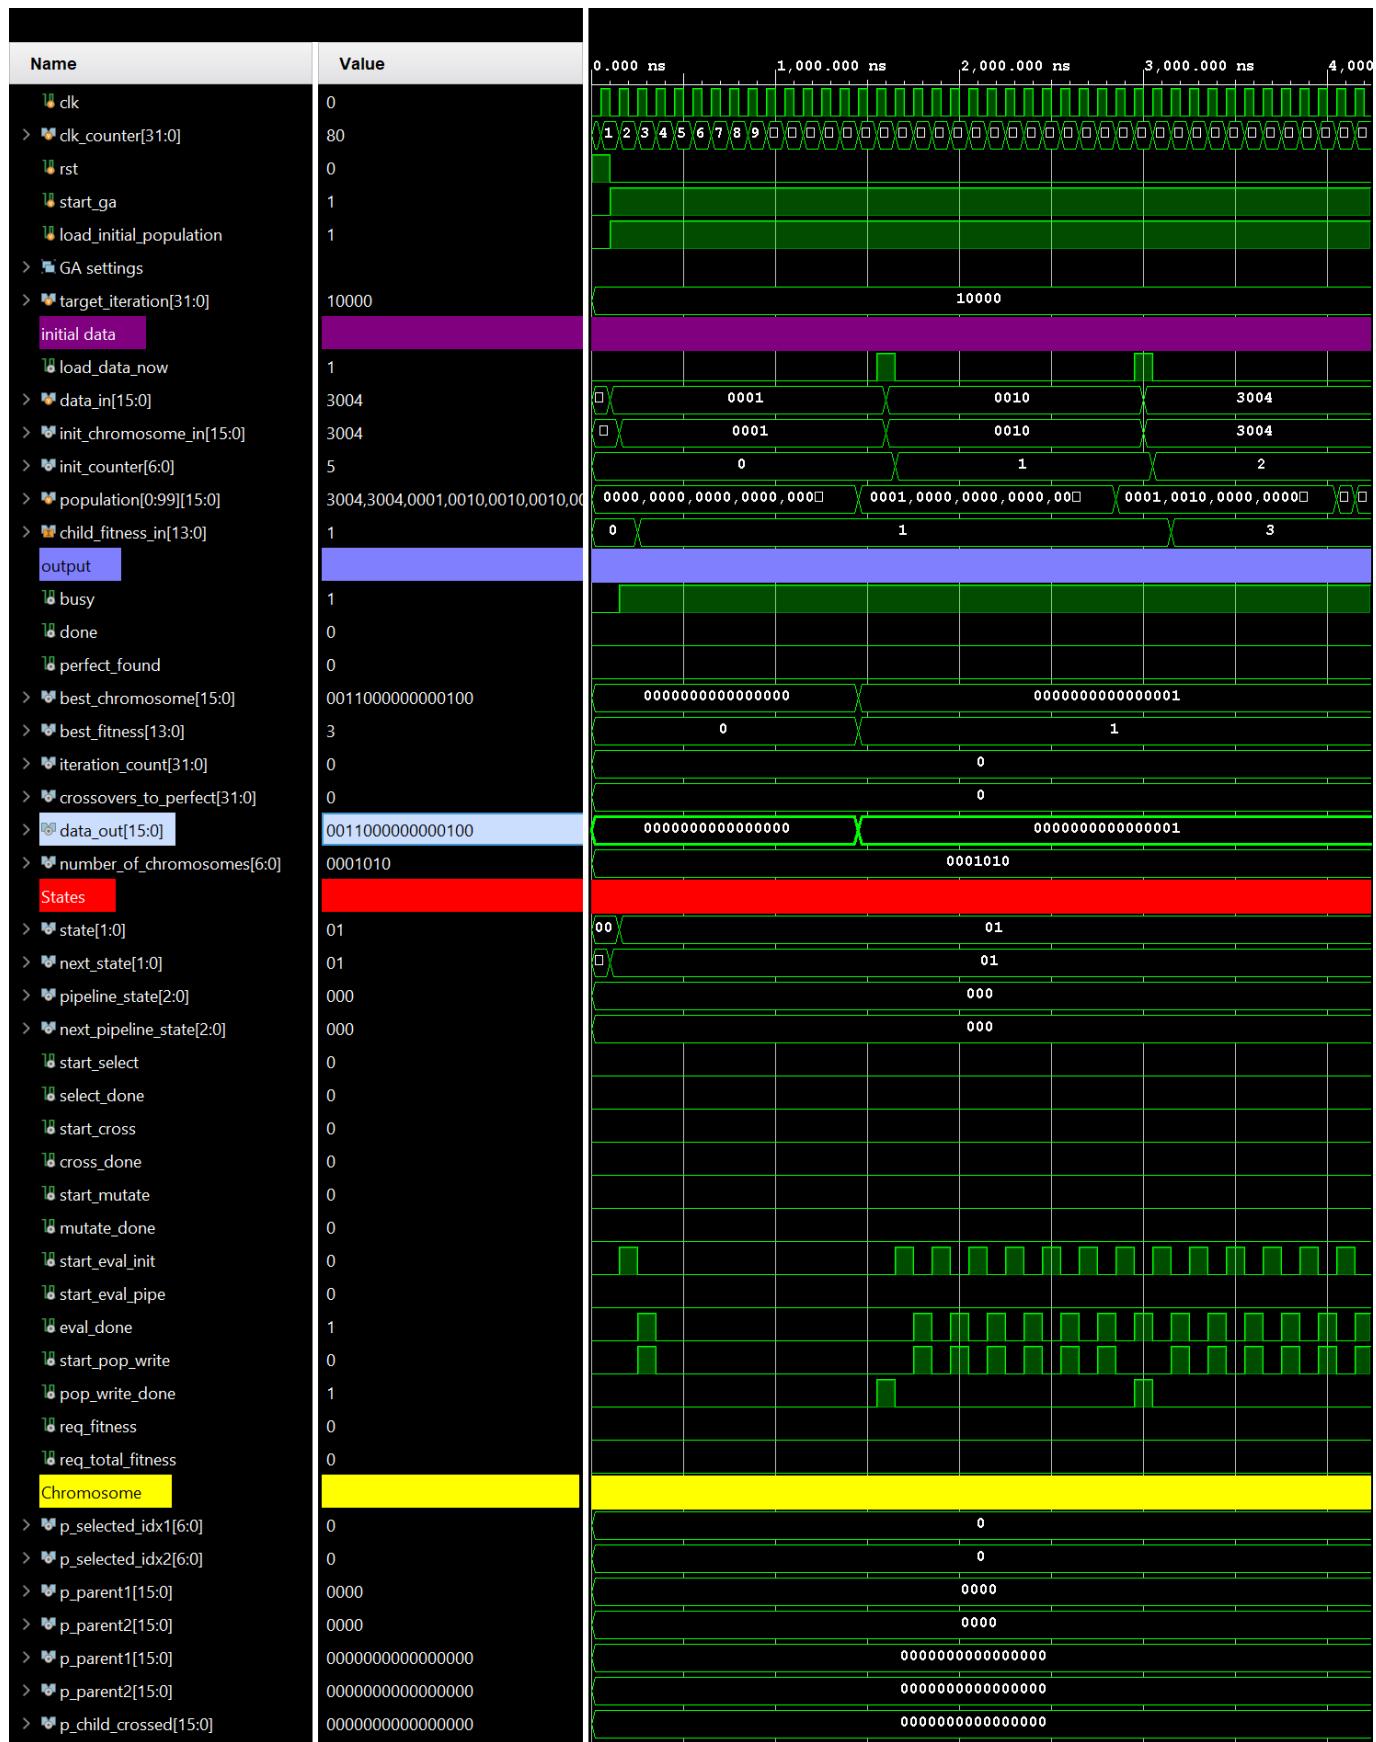
simulation:

Name	Value	706,400.000 ns	706,600.000 ns	706,800.000 ns	707,000.000 ns	707,200.000 ns	707,400.000 ns
		707,550,000 ns					

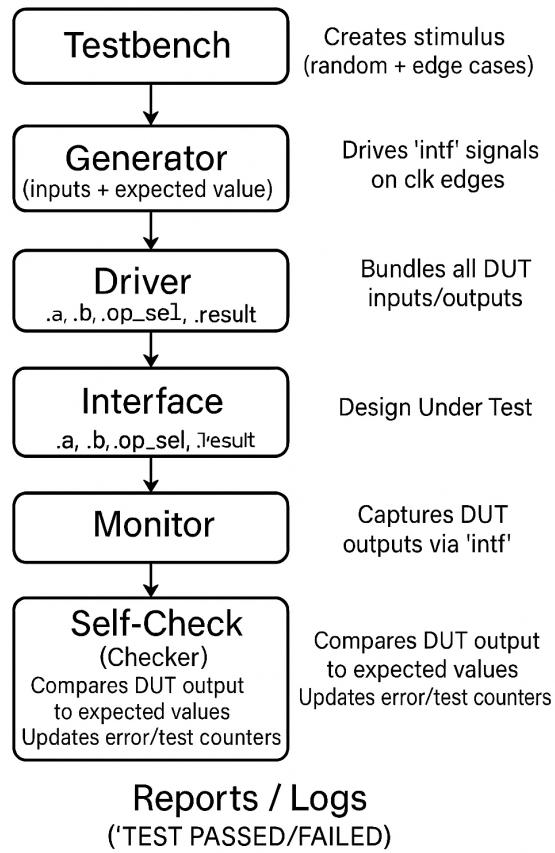




delay B-ram is much higher then LUT:



TestBench style



Reports / Logs
('TEST PASSED/FAILED')

```

`timescale 1ns / 1ps
// -----
// 1. Interface – bundles DUT I/O signals
// -----
interface dut_if(input logic clk, rst);
    // DUT inputs
    logic [7:0] a;
    logic [7:0] b;
    logic      carry_in;
    logic [1:0] op_sel;

    // DUT outputs
    logic [7:0] result;
    logic      carry_out;
endinterface

// -----
// 2. Testbench module
// -----
module tb_top;

    // Clock generation (100 MHz with timescale 1ns/1ps)
    logic clk = 0;
    always #5 clk = ~clk;

    // Reset generation
    logic rst;
    initial begin
        rst = 1;
        #20 rst = 0;
    end

    // Interface instance
    dut_if intf(clk, rst);

    // Expected values (based on reference / golden model)
    logic [7:0] expected_result;
    logic      expected_carry;

    int error_count = 0;
    int test_count  = 0;

    // -----
    // 3. DUT instantiation
    // -----
    // Replace 'my_dut' and port connections as per your DUT
    my_dut DUT (
        .a(intf.a),
        .b(intf.b),
        .carry_in(intf.carry_in),

```

```

    .op_sel(intf.op_sel),
    .result(intf.result),
    .carry_out(intf.carry_out)
);

// -----
// 4. Generator – produces random + edge case stimuli
// -----
task generator(int num_tests);
    for (int i = 0; i < num_tests; i++) begin
        // Default randomization
        intf.a      = $urandom();
        intf.b      = $urandom();
        intf.carry_in = $urandom_range(0, 1);
        intf.op_sel  = $urandom_range(0, 3);

        // Inject example edge cases at fixed intervals
        if (i % 100 == 0) begin
            intf.a = 8'h00;
            intf.b = 8'hFF;
        end

        // Compute expected values based on your DUT function
        // (Replace this block with your golden model logic)
        case (intf.op_sel)
            2'b00: {expected_carry, expected_result} = intf.a + intf.b;
            2'b01: {expected_carry, expected_result} = {1'b0, intf.a} - {1'b0,
intf.b};
            2'b10: begin expected_result = intf.a & intf.b; expected_carry = 0;
        end
            2'b11: begin expected_result = intf.a | intf.b; expected_carry = 0;
        end
        endcase

        // Send to DUT via driver
        driver();
    end
endtask

// -----
// 5. Driver – applies data to DUT synchronised to clock
// -----
task driver();
    @(posedge clk);
    #1; // allow small time for signals to settle
    monitor();
    test_count++;
endtask

// -----
// 6. Monitor – captures DUT outputs and calls checker

```

```

// -----
task monitor();
    check_result();
endtask

// -----
// 7. Checker – compares DUT output with expected values
// -----
task check_result();
    if (intf.result !== expected_result) begin
        $display("ERROR @ %0t: op_sel=%b, a=%h, b=%h | expected result=%h",
got=%h",
                $time, intf.op_sel, intf.a, intf.b,
                expected_result, intf.result);
        error_count++;
    end
    if (intf.carry_out !== expected_carry) begin
        $display("ERROR @ %0t: op_sel=%b, a=%h, b=%h | expected carry=%b, got=%b",
                $time, intf.op_sel, intf.a, intf.b,
                expected_carry, intf.carry_out);
        error_count++;
    end
endtask

// -----
// 8. Main sequence – controls simulation flow
// -----
initial begin
    wait(!rst); // Wait until reset is released
    $display("Starting tests...");
    generator(1000); // Number of tests
    $display("Tests completed: %0d total, %0d errors", test_count, error_count);

    if (error_count == 0)
        $display("TEST PASSED");
    else
        $display("TEST FAILED");

    $finish;
end

// Optional waveform dump
initial begin
    $dumpfile("tb_top.vcd");
    $dumpvars(0, tb_top);
end

endmodule

```